



THE UNIVERSITY OF  
CHICAGO

**MASTERS** IN  
COMPUTATIONAL  
SOCIAL SCIENCE  
THE UNIVERSITY OF CHICAGO

A large, semi-transparent image of Po the panda from the movie Kung Fu Panda. He is wearing his signature conical straw hat and a red and yellow scarf. He has a confident, slightly smug expression. The background of the image is a soft-focus view of a traditional Chinese village with tiled roofs.

**MACS 30111**

**M7: Pandas**

# Agenda / misc

- ▶ With Pandas, easiest option is jupyter notebook

# Topics:

- ❑ **Introduction to Pandas**
- ❑ Creating DataFrames and Series
- ❑ Working with DataFrames and Series
- ❑ Applying functions to data
- ❑ Group, combine, and pivoting
- ❑ Visualizing DataFrames and Series (Probably next week)



Data analysis toolkit

Reference: <https://pandas.pydata.org/>

# The “*pandas*” name

*pan* for “Panel”

*da* for “Data”

Panel data: multidimensional, structured datasets that include observations of something over a number of different time periods.

# Use cases for the pandas library

Tabular Data

Time-series Data

Matrix Data

Statistical Datasets

# Topics:

- ❑ Introduction to Pandas
- ❑ **Creating DataFrames and Series**
- ❑ Working with DataFrames and Series
- ❑ Applying functions to data
- ❑ Group, combine, and pivoting
- ❑ Visualizing DataFrames and Series

# Pandas

## Series

In: s

Out:

```
Rory    90
Lorelai  95
Luke     90
dtype: int64
```

## DataFrame

In: s = pd.DataFrame(data)

In: s

Out:

```
      PA1  PA2
Rory    90   88
Lorelai  95   90
Luke     90   75
```

### Reference:

- <https://pandas.pydata.org/docs/reference/api/pandas.Series.html>
- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>



# Importing the Pandas library

Import pandas using the alias “pd”

```
import pandas as pd
```

# Series Constructor

Pass in 1-d data (list/np.array), alongside an accompanying (1-d) index:

```
data = [90, 95, 90]
```

```
index = ["Rory", "Lorelai", "Luke"]
```

```
s = pd.Series(data, index)
```

s

Out:

Rory 90

Lorelai 95

Luke 90

dtype: int64

Pass in a dictionary, where the keys are inferred to be the indices:

```
data = {"Rory":90, "Lorelai": 95, "Luke": 90}
```

```
s = pd.Series(data)
```

# DataFrame Constructor

Pass in a dictionary of lists/arrays, alongside an accompanying (1-d) index:

```
data = {"PA1": [90,95,90], "PA2": [88,90,75]}
```

```
index = ["Rory", "Lorelai", "Luke"]
```

```
df = pd.DataFrame(data, index = index)
```

|         | PA1 | PA2 |
|---------|-----|-----|
| Rory    | 90  | 88  |
| Lorelai | 95  | 90  |
| Luke    | 90  | 75  |

Pass in a **dictionary of dictionaries**

```
data = {"PA1": {"Rory": 90, "Lorelai": 95, "Luke": 90},  
        "PA2": {"Rory": 88, "Lorelai": 90, "Luke": 75}}
```

# DataFrame Constructor

Data can also be passed in as an NumPy multidimensional array:

```
In: import numpy as np
```

```
In: np.random.seed(13)
```

```
In: data = np.random.uniform(0, 1, (3,2))
```

```
In: data
```

```
Out:
```

```
array([[0.77770241, 0.23754122],  
       [0.82427853, 0.9657492 ],  
       [0.97260111, 0.45344925]])
```

# DataFrame Constructor

Data can also be passed in as an NumPy multidimensional array:

```
In: index = ["Rory", "Lorelai", "Luke"]
```

```
In: columns = ["pa1", "pa2"]
```

```
In: df = pd.DataFrame(data, index, columns)
```

```
In: df
```

```
Out:
```

|         | pa1      | pa2      |
|---------|----------|----------|
| Rory    | 0.777702 | 0.237541 |
| Lorelai | 0.824279 | 0.965749 |
| Luke    | 0.972601 | 0.453449 |

# Read from file

CSV

JSON

HTML

Excel Spreadsheets

And much more...

```
trees = pd.read_csv("2015_Street_Tree_Census_Tree_Data_20231117.csv")
```

Reference: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/io.html](https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html)

# Key commands

- ↳ Loading: `read_csv()`
- ↳ Summarizing:
  - ↳ `data.head()`
  - ↳ `data.tail()`
- ↳ Selecting:
  - ↳ By index: `iloc` (e.x. `trees.iloc[:,1]`)
  - ↳ By name: `loc` (e.x. `trees.loc[:, "block_id"]`)

# Series VS DataFrame

## ► Series:

- One-dimensional ndarray with axis labels
- |         |    |
|---------|----|
| Rory    | 90 |
| Lorelai | 95 |
| Luke    | 90 |

## ► DataFrame

- Two-dimensional tabular data with labeled row and column axes

|         | PA1 | PA2 |
|---------|-----|-----|
| Rory    | 90  | 88  |
| Lorelai | 95  | 90  |
| Luke    | 90  | 75  |

## Reference:

- Series: <https://pandas.pydata.org/docs/reference/api/pandas.Series.html>
- DataFrame: <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>



# Approaching the Series / DataFrame

Size, column names and data types: *dtype*, *shape*

```
In [34]: s
```

```
Out[34]:
```

```
Rory    90
Lorelai 95
Luke    90
dtype: int64
```

```
In [35]: s.dtype
```

```
Out[35]: dtype('int64')
```

```
In [36]: s.shape
```

```
Out[36]: (3,)
```

```
In [37]: df
```

```
Out[37]:
```

```
      PA1 PA2
Rory    90  88
Lorelai 95  90
Luke    90  75
```

```
In [38]: df.dtypes
```

```
Out[38]:
```

```
PA1    int64
PA2    int64
dtype: object
```

```
In [39]: df.shape
```

```
Out[39]: (3, 2)
```

# Approaching the Series / DataFrame

What does the data look like? *df.head()*, *df.tail()*

```
In [40]: s.head(n=2)
```

```
Out[40]:
```

```
Rory    90
```

```
Lorelai 95
```

```
dtype: int64
```

```
In [41]: s.tail(n=1)
```

```
Out[41]:
```

```
Luke    90
```

```
dtype: int64
```

# Missing Data

Identifying missing data in Series and DataFrames:

```
s.isna(); s.notna()  
df.isna(); df.notna()
```

We can drop rows or columns with missing data:

```
df_drop_byrow = df.dropna()  
df_drop_bycol = df.dropna(axis = 1)
```

We can also fill missing values:

```
df.col.fillna(np.mean(df.col))
```

```
In [42]: s.isna()  
Out[42]:  
Rory    False  
Lorelai False  
Luke    False  
dtype: bool
```

```
In [43]: s.notna()  
Out[43]:  
Rory    True  
Lorelai True  
Luke    True  
dtype: bool
```

```
In [44]: df.isna()  
Out[44]:  
      PA1  PA2  
Rory  False False  
Lorelai False False  
Luke  False False
```

```
In [45]: df.notna()  
Out[45]:  
      PA1  PA2  
Rory   True  True  
Lorelai True  True  
Luke   True  True
```

Coding practice: 4.3.2

# Motivating example

- ▶ Data exploration for the 2015 New York City Street Tree Survey
- ▶ Data source: <https://data.cityofnewyork.us/Environment/2015-Street-Tree-Census-Tree-Data/uvpi-gqnh>
- ▶ Go to box: <https://uchicago.box.com/s/t4ix9us983xsq3x9lqhpsmnzeogte349>

1. How many different species are planted as street trees in New York?
2. What are the five most common street tree species in New York?
3. What is the most common street tree species in Brooklyn?
4. What percentage of the street trees in Queens are dead or in poor health?
5. How does street tree health differ by borough?

[https://github.com/computer-science-with-applications/examples/tree/main/working\\_with\\_data/pandas](https://github.com/computer-science-with-applications/examples/tree/main/working_with_data/pandas)

Coding practice: 4.3.1

# Selecting rows (indexing)

```
In [6]: trees.head()
```

```
Out[6]:
```

|   | tree_id | block_id | created_at | tree_dbh | ... | council district | census tract    | bin          | bbl |
|---|---------|----------|------------|----------|-----|------------------|-----------------|--------------|-----|
| 0 | 180683  | 348711   | 08/27/2015 | 3        | ... | 29.0             | 739.0 4052307.0 | 4.022210e+09 |     |
| 1 | 200540  | 315986   | 09/03/2015 | 21       | ... | 19.0             | 973.0 4101931.0 | 4.044750e+09 |     |
| 2 | 204026  | 218365   | 09/05/2015 | 3        | ... | 34.0             | 449.0 3338310.0 | 3.028870e+09 |     |
| 3 | 204337  | 217969   | 09/05/2015 | 10       | ... | 34.0             | 449.0 3338342.0 | 3.029250e+09 |     |
| 4 | 189565  | 223043   | 08/30/2015 | 21       | ... | 39.0             | 165.0 3025654.0 | 3.010850e+09 |     |

## Row label-based indexing

```
In [7]: trees.loc[180683]
```

```
Out[7]:
```

|            |            |
|------------|------------|
| tree_id    | 362889     |
| block_id   | 209521     |
| created_at | 10/22/2015 |
| tree_dbh   | 6          |
| ...        |            |

## Position-based indexing

```
In [8]: trees.iloc[0]
```

```
Out[8]:
```

|            |             |
|------------|-------------|
| tree_id    | 180683      |
| block_id   | 348711      |
| created_at | 08/27/2015  |
| tree_dbh   | 3           |
| stump_diam | 0           |
| curb_loc   | OnCurb      |
| status     | Alive       |
| health     | Fair        |
| spc_latin  | Acer rubrum |
| ...        |             |

# Naming columns

↓ If you want to explore the column names that were imported, you can use: `trees.columns`

```
Index(['tree_id', 'block_id', 'created_at', 'tree_dbh', 'stump_diam',  
      'curb_loc', 'status', 'health', 'spc_latin', 'spc_common', 'steward',  
      'guards', 'sidewalk', 'user_type', 'problems', 'root_stone',  
      'root_grate', 'root_other', 'trunk_wire', 'trnk_light', 'trnk_other',  
      'brch_light', 'brch_shoe', 'brch_other', 'address', 'postcode',  
      'zip_city', 'community board', 'borocode', 'borough', 'cnclldist',  
      'st_assem', 'st_senate', 'nta', 'nta_name', 'boro_ct', 'state',  
      'latitude', 'longitude', 'x_sp', 'y_sp', 'council district',  
      'census tract', 'bin', 'bbl'],  
      dtype='object')
```

# Selecting columns

*Note: depending on data, might have 'boroname'*

```
1 type(trees['health'])
```

pandas.core.series.Series

```
In [16]: trees["health"]
```

```
Out[16]:
```

|   |      |
|---|------|
| 0 | Fair |
| 1 | Fair |
| 2 | Good |
| 3 | Good |
| 4 | Good |

...

|        |      |
|--------|------|
| 683783 | Good |
| 683784 | Good |
| 683785 | Good |
| 683786 | Good |
| 683787 | Fair |

Name: health, Length: 683788, dtype: object

```
1 type(trees[['health']])
```

pandas.core.frame.DataFrame

```
In [17]: trees[["health"]]
```

```
Out[17]:
```

|   | health |
|---|--------|
| 0 | Fair   |
| 1 | Fair   |
| 2 | Good   |
| 3 | Good   |
| 4 | Good   |

...

|        |      |
|--------|------|
| 683783 | Good |
| 683784 | Good |
| 683785 | Good |
| 683786 | Good |
| 683787 | Fair |

[683788 rows x 1 columns]

```
1 type(trees[['health', 'borough']])
```

pandas.core.frame.DataFrame

```
In [18]: trees[["health", "borough"]]
```

```
Out[18]:
```

|        | health | borough       |
|--------|--------|---------------|
| 0      | Fair   | Queens        |
| 1      | Fair   | Queens        |
| 2      | Good   | Brooklyn      |
| 3      | Good   | Brooklyn      |
| 4      | Good   | Brooklyn      |
| ...    | ...    | ...           |
| 683783 | Good   | Brooklyn      |
| 683784 | Good   | Queens        |
| 683785 | Good   | Staten Island |
| 683786 | Good   | Bronx         |
| 683787 | Fair   | Queens        |

[683788 rows x 2 columns]

# Selecting rows and columns

*Note: depending on data, might have 'boroname'*

```
In [19]: trees.iloc[:5][["health", "borough"]]
```

```
Out[19]:
```

|   | health | borough  |
|---|--------|----------|
| 0 | Fair   | Queens   |
| 1 | Fair   | Queens   |
| 2 | Good   | Brooklyn |
| 3 | Good   | Brooklyn |
| 4 | Good   | Brooklyn |

```
In [20]: trees.loc[[180683, 200540, 204026], ["health",  
...: "borough"]]
```

```
Out[20]:
```

|        | health | borough       |
|--------|--------|---------------|
| 180683 | Good   | Brooklyn      |
| 200540 | NaN    | Staten Island |
| 204026 | Good   | Staten Island |



# Filtering

- Boolean masks
- Logical operations

```
filter = (trees.borough == "Queens") & ((trees.status == "Dead") | (trees.health == "Poor"))  
trees[filter]
```

- Remove rows where trees.status is Dead
- Select rows where trees.status is not Dead

```
trees[trees.status != "Dead"]
```

```
0    False  
1    False  
2    False  
3    False  
4    False
```

```
...  
683783  False  
683784  False  
683785  False  
683786  False  
683787  False
```

```
Name: status, Length: 683788, dtype:  
bool
```

# Agenda / misc

- ▶ Applied practice
- ▶ Room reservation: final exam
- ▶ Deadlines/grading/etc.
- ▶ Anything else?

# Topics:

- ❑ Introduction to Pandas
- ❑ Creating DataFrames and Series
- ❑ **Working with DataFrames and Series**
- ❑ Applying functions to data
- ❑ Group, combine, and pivoting
- ❑ Visualizing DataFrames and Series (Probably Thurs)

# Review: load in data + add row names

- ↳ How to load data?
- ↳ How to set row names? (and why)
  - ↳ Hint: `index_col`
- ↳ **Application:** find tree at index 9 and tree\_id of 192755

# Practice Task: answer a research question

- ↳ **Suppose we have the following question:** how do tree health indicators vary by borough? **\*\*GIVEN THE CURRENT SKILLS / TOOLS WE HAVE DISCUSSED SO FAR\*\***
- ↳ **In groups:**
  - ↳ Determine how you will measure this. (which column)
  - ↳ What tools do you currently have to do this?
  - ↳ What do you want to be able to do?

# Topics:

- ❑ Introduction to Pandas
- ❑ Creating DataFrames and Series
- ❑ Working with DataFrames and Series
- ❑ **Applying functions to data**
- ❑ Group, combine, and pivoting
- ❑ Visualizing DataFrames and Series

# Differences: help me understand

```
In : trees.count()
```

```
Out :
```

|            |        |
|------------|--------|
| tree_id    | 683788 |
| block_id   | 683788 |
| created_at | 683788 |
| tree_dbh   | 683788 |
| stump_diam | 683788 |
| curb_loc   | 683788 |
| status     | 683788 |
| health     | 652172 |
| spc_latin  | 652169 |
| spc_common | 652169 |

```
In : trees.status.value_counts()
```

```
Out :
```

|        |        |
|--------|--------|
| status |        |
| Alive  | 652173 |
| Stump  | 17654  |
| Dead   | 13961  |

Name: count, dtype: int64

```
In : trees.count()
```

```
Out :
```

|            |        |
|------------|--------|
| block_id   | 683788 |
| created_at | 683788 |
| tree_dbh   | 683788 |
| stump_diam | 683788 |
| curb_loc   | 683788 |
| status     | 683788 |
| health     | 652172 |
| spc_latin  | 652169 |
| spc_common | 652169 |

```
In : trees.status.value_counts()
```

```
Out:
```

|        |        |
|--------|--------|
| status |        |
| Alive  | 652173 |
| Stump  | 17654  |
| Dead   | 13961  |

Name: count, dtype: int64

Reference: <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.describe.html>

# Common Summary Statistics

## `DataFrame.count`

Count number of non-NA/null observations.

## `DataFrame.max`

Maximum of the values in the object.

## `DataFrame.min`

Minimum of the values in the object.

## `DataFrame.mean`

Mean of the values.

## `DataFrame.std`

Standard deviation of the observations.

## `DataFrame.select_dtypes`

Subset of a DataFrame including/excluding columns based on their dtype.

Summary statistics for each column: `df.describe()`

Summary of categorical values: `df.value_counts()`

Pairwise correlation of columns: `df.corr()`

```
In : trees.count()
```

```
Out :
```

```
block_id      683788
created_at    683788
tree_dbh      683788
stump_diam     683788
curb_loc      683788
status        683788
health        652172
spc_latin     652169
spc_common    652169
```

```
In : trees.status.value_counts()
```

```
Out :
```

```
status
Alive  652173
Stump  17654
Dead   13961
Name: count, dtype: int64
```

Reference: <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.describe.html>



# Sorting

- ↓ `df.sort_values()`
- ↓ `df.describe()`

# Applying Functions to Data

## `Series.apply`

For applying more complex functions on a Series.

## `DataFrame.apply`

Apply a function row-/column-wise.

## `DataFrame.applymap`

Apply a function elementwise on a whole DataFrame.

# Applying functions: example

```
def from_celsius(temp):
```

```
    return temp*9/5 + 32
```

```
import pandas as pd
```

```
df = pd.DataFrame([list(range(0,100,10)), list(range(5,105, 10))])
```

```
df.apply(from_celsius)
```

# Topics:

- ❑ Introduction to Pandas
- ❑ Creating DataFrames and Series
- ❑ Working with DataFrames and Series
- ❑ Applying functions to data
- ❑ **Group, combine, and pivoting**
- ❑ Visualizing DataFrames and Series

# Grouping

```
In: trees.groupby("borough")[["tree_dbh"]].mean()
```

Out:

|               | tree_dbh  |
|---------------|-----------|
| borough       |           |
| Bronx         | 9.693649  |
| Brooklyn      | 11.738884 |
| Manhattan     | 8.473641  |
| Queens        | 12.557870 |
| Staten Island | 10.492746 |

# Grouping

```
In : trees.groupby(["borough", "status"]).size()
```

```
Out :
```

| borough       | status |        |
|---------------|--------|--------|
| Bronx         | Alive  | 80585  |
|               | Dead   | 2530   |
|               | Stump  | 2088   |
| Brooklyn      | Alive  | 169744 |
|               | Dead   | 3319   |
|               | Stump  | 4230   |
| Manhattan     | Alive  | 62427  |
|               | Dead   | 1802   |
|               | Stump  | 1194   |
| Queens        | Alive  | 237974 |
|               | Dead   | 4440   |
|               | Stump  | 8137   |
| Staten Island | Alive  | 101443 |
|               | Dead   | 1870   |
|               | Stump  | 2005   |

```
dtype: int64
```

# Reshaping: Stacking / Unstacking

- ▶ Think about maneuvering your data around
  - what do you want your rows and columns to be - what is the observation?
- ▶ Do you want an indicator variable, a count variable, etc?

Unstack: converts series with hierarchical index into a DataFrame

# Unstacking

By default, unstack will order the columns by value

```
In : trees.groupby(["borough", "status"]).size()
```

Out :

| borough  | status |        |
|----------|--------|--------|
| Bronx    | Alive  | 80585  |
|          | Dead   | 2530   |
|          | Stump  | 2088   |
| Brooklyn | Alive  | 169744 |
| ...      |        |        |

```
In : trees.groupby(["borough", "status"]).size().unstack()
```

Out:

| status        | Alive  | Dead | Stump |
|---------------|--------|------|-------|
| borough       |        |      |       |
| Bronx         | 80585  | 2530 | 2088  |
| Brooklyn      | 169744 | 3319 | 4230  |
| Manhattan     | 62427  | 1802 | 1194  |
| Queens        | 237974 | 4440 | 8137  |
| Staten Island | 101443 | 1870 | 2005  |



# Good exam question: Programming practice

- ▶ CLOSED NOTE
- ▶ Suppose you have a dataframe with 20 rows and 10 columns.
  - ▶ DF name: albums
  - ▶ Column names: title, num\_songs, year, genre, length, billboard, weeks\_on\_charts, label, avg\_bpm, number\_sold
- ▶ From the above dataframe, write the following code on your paper:
  - ▶ Get the shape of the dataframe
  - ▶ Print the information for the album in the 10<sup>th</sup> row
  - ▶ Print the information for the “genre” column
  - ▶ (bonus) Select albums that have have been in the top 100 on billboard (e.g. have a value less than or equal to 100)

# TASK: develop a research question

- ↴ In your groups, develop a research question you can answer with the trees dataset.
- ↴ Your question needs to use ALL of the following:
  - ↵ Filtering
  - ↵ Applying a Boolean criterion
  - ↵ Label-based indexing
  - ↵ Calculating a numerical value (e.g. mean/min/max/std/corr)
  - ↵ Exporting the text file (documentation!)

# Using functions

What does this do/mean? Why and how would we do this?

```
for col_name in ["boroname", "health", "spc_common", "status"]:  
    trees[col_name] = trees[col_name].astype("category")
```

# Using functions

↳ What does this do/mean? Why and how would we do this?

```
def tree_health_by_boro(trees):  
    combined_status = trees.status.where(trees.status != "Alive", trees.health)  
    num_per_boro = trees.groupby("borough").size()  
    combined_per_boro = trees.groupby(["borough", combined_status]).size()  
    pct_per_boro = combined_per_boro/num_per_boro*100.0  
    pct_per_boro_df = pct_per_boro.unstack()  
  
    return pct_per_boro_df[["Good", "Fair", "Poor", "Dead", "Stump"]]
```

```
tree_health_by_boro(trees)
```

# Unstacking

By default, unstack will order the columns by value

```
In : trees.groupby(["borough", "status"]).size()
```

Out :

| borough  | status |        |
|----------|--------|--------|
| Bronx    | Alive  | 80585  |
|          | Dead   | 2530   |
|          | Stump  | 2088   |
| Brooklyn | Alive  | 169744 |
| ...      |        |        |

```
In : trees.groupby(["borough", "status"]).size().unstack()
```


Out:

| status        | Alive  | Dead | Stump |
|---------------|--------|------|-------|
| borough       |        |      |       |
| Bronx         | 80585  | 2530 | 2088  |
| Brooklyn      | 169744 | 3319 | 4230  |
| Manhattan     | 62427  | 1802 | 1194  |
| Queens        | 237974 | 4440 | 8137  |
| Staten Island | 101443 | 1870 | 2005  |

**“Happy families are all alike; every unhappy family is unhappy in its own way.” -- Leo Tolstoy**


**“Tidy datasets are all alike, but every messy dataset is messy in its own way.” -- Hadley Wickham**

The standard structure of tidy data means that  
"tidy datasets are all alike..."



"...but every messy dataset is  
messy in its own way."

—HADLEY WICKHAM



“**TIDY DATA** is a standard way of mapping the meaning of a dataset to its structure.”

—HADLEY WICKHAM

## In tidy data:

- each variable forms a column
- each observation forms a row
- each cell is a single measurement

each column a variable



| id | name   | color  |
|----|--------|--------|
| 1  | floof  | gray   |
| 2  | max    | black  |
| 3  | cat    | orange |
| 4  | donut  | gray   |
| 5  | merlin | black  |
| 6  | panda  | calico |

each row  
an  
observation





# Long vs wide

wide

| id | x | y | z |
|----|---|---|---|
| 1  | a | c | e |
| 2  | b | d | f |

long

| id | key | val |
|----|-----|-----|
| 1  | x   | a   |
| 2  | x   | b   |
| 1  | y   | c   |
| 2  | y   | d   |
| 1  | z   | e   |
| 2  | z   | f   |

## Benefits: long vs wide

‘Long’ data: better for analysis

‘Wide’ data: can be helpful for visualization

# Long vs wide

| Common Name     | lifespan_cat_eq_Short | lifespan_cat_eq_Medium | lifespan_cat_eq_Long | lifespan_cat_eq_ExtraLong | Common Name | lifespan_cat              | value |
|-----------------|-----------------------|------------------------|----------------------|---------------------------|-------------|---------------------------|-------|
| Ash, Green      | 0                     | 1                      | 0                    | 0                         | Ash, Green  | lifespan_cat_eq_Short     | 0     |
| Ash, White      | 0                     | 0                      | 0                    | 1                         | Ash, Green  | lifespan_cat_eq_Medium    | 1     |
| Basswood        | 1                     | 0                      | 0                    | 0                         | Ash, Green  | lifespan_cat_eq_ExtraLong | 0     |
| Beech, American | 0                     | 0                      | 0                    | 1                         | Ash, Green  | lifespan_cat_eq_Long      | 0     |
| Birch, Gray     | 1                     | 0                      | 0                    | 0                         | Ash, White  | lifespan_cat_eq_ExtraLong | 1     |

# Pivoting

Convert a long and thin DataFrame into a short and wide DataFrame

- A column for index
- A column to supply column names for the new DataFrame
- A column to fill the new DataFrame

**Options:** pivot, melt, wide\_to\_long

**Coding practice: 4.3.6**

# Pivot example

```
pd.melt(trees_life, value_vars=filter_col,  
id_vars= ["Common Name", "Scientific Name"],  
var_name='lifespan_cat')
```

```
pd.wide_to_long(trees_life, stubnames="lifespan_cat_eq",  
i="Common Name",  
j="lifespan_cat", suffix='\\w+')
```

```
trees_life.pivot(index='Common Name', columns='lifespan_cat', values='Average_lifespan')
```

```
(bonus) pd.get_dummies(trees_life, columns=['lifespan_cat_eq'], dtype=int)
```

# Combining DataFrames

- **concat()**
  - perform concatenation along an axis
  - while performing set logic of the indexes on other axes
  - make a full copy of the data
- **merge()**
  - Standard database join operations between DataFrame or Series objects
- **join()**
  - join on index
  - combine columns of two differently indexed DataFrames into a single one

**Reference:** [https://pandas.pydata.org/docs/user\\_guide/merging.html](https://pandas.pydata.org/docs/user_guide/merging.html)

# Join vs merge

- ↳ Merge is the ‘big picture’ for things
  - ↳ Can have more freedom with indices / how you merge
- ↳ ‘join’ is like a subset of merge --

# Merge: syntax

↳ **`DataFrame.merge(right, how='inner', on=None, left_on=None, right_on=None, left_index=False, right_index=False, sort=False, suffixes=('_x', '_y'), copy=None, indicator=False, validate=None)`**

## Parameters:

- **how**{'left', 'right', 'outer', 'inner', 'cross'}, default 'inner' Type of merge to be performed.
  - left: use only keys from left frame, similar to a SQL left outer join; preserve key order.
  - right: use only keys from right frame, similar to a SQL right outer join; preserve key order.
  - outer: use union of keys from both frames, similar to a SQL full outer join; sort keys lexicographically.
  - inner: use intersection of keys from both frames, similar to a SQL inner join; preserve the order of the left keys.
  - cross: creates the cartesian product from both frames, preserves the order of the left keys.
- **On** label or listColumn or index level names to join on. These must be found in both DataFrames. If *on* is None and not merging on indexes then this defaults to the intersection of the columns in both DataFrames.



# Merge options

1. `df1.merge(df2, how = "right")`
2. `df1.merge(df2, how = "left")`
3. `df1.merge(df2, how = "inner")`
4. `df1.merge(df2, how = "outer")`

# Syntax: is it `df.merge` or `pd.merge`?

Similar across circumstances but just depends on how you want to call it...when you look at the documentation, you'll see this trend consistently.

If you call `pd.function`, you will need to specify the data frame(s).

If you `df.function`, you no longer need to specify the data frame.

# TASK: ADVENTURE TIME (cont'd)

1. Get to know your second dataset
2. Merge the dataframes
3. Think about a question you could ask and give it a go

# Announcements

## ▶ **Final exam schedule**

- ▶ Tues, Dec 9<sup>h</sup>, 12:30 PM-2:30 PM
- ▶ Take-home: open Friday through Weds
  - ▶ Timing: 36 hours (same as midterm)

## ▶ **Review**

- ▶ Textbook
- ▶ Team tutorial
- ▶ Short exercise
- ▶ Programming assignment
- ▶ Extra exercise